

Made by Ansh Sharma- Sentimental Analysis by Hugging Face Fine-Tuning BERT , Deployed on wandb AI, Heroku through FASTAPI and Local Server, followed by Testing and Evaluating + Productization of the Model

Attached with various Screenshots , and server Analysis.

```
1 pip install transformers datasets scikit-learn torch matplotlib
```



Show hidden output

Preparing Synthetic Data and Dividing dataset into Training and Testing

```
1 from datasets import Dataset
2 from sklearn.model_selection import train_test_split
3 import pandas as pd
4
5 # Example synthetic dataset
6 data = {
7     "ReviewText": [
8         "The event was amazing and well-organized!",
9         "The amenities were not up to the mark.",
10        "Had a decent experience overall.",
11        "The campus event was a complete disaster.",
12        "I loved every moment of the event!"
13    ],
14    "SentimentLabel": [2, 0, 1, 0, 2] # 0: Negative, 1: Neutral, 2: Positive
15 }
16
17 # Create a DataFrame
18 df = pd.DataFrame(data)
19
20 # Split into train, validation, and test sets
21 train_texts, test_texts, train_labels, test_labels = train_test_split(
22     df["ReviewText"], df["SentimentLabel"], test_size=0.2, random_state=42
23 )
24
25 train_texts, val_texts, train_labels, val_labels = train_test_split(
26     train_texts, train_labels, test_size=0.1, random_state=42
27 )
28
29 # Convert to Hugging Face Dataset
```

```

30 train_dataset = Dataset.from_dict({"text": train_texts, "label": train_labels})
31 val_dataset = Dataset.from_dict({"text": val_texts, "label": val_labels})
32 test_dataset = Dataset.from_dict({"text": test_texts, "label": test_labels})
33

```

Training and Tokenize - Hugging Face , Wandb AI-And API Key

```

1 from transformers import BertTokenizer, BertForSequenceClassification, Trainer, TrainingArguments
2 import torch
3
4 # Load BERT tokenizer and model
5 tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
6 model = BertForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=3)
7
8 # Tokenize datasets
9 def tokenize_function(examples):
10     return tokenizer(examples["text"], padding="max_length", truncation=True)
11
12 train_dataset = train_dataset.map(tokenize_function, batched=True)
13 val_dataset = val_dataset.map(tokenize_function, batched=True)
14 test_dataset = test_dataset.map(tokenize_function, batched=True)
15
16 # Set format for PyTorch
17 train_dataset.set_format("torch", columns=["input_ids", "attention_mask", "label"])
18 val_dataset.set_format("torch", columns=["input_ids", "attention_mask", "label"])
19 test_dataset.set_format("torch", columns=["input_ids", "attention_mask", "label"])
20
21 # Training arguments
22 training_args = TrainingArguments(
23     output_dir="./results",
24     evaluation_strategy="epoch",
25     learning_rate=2e-5,
26     per_device_train_batch_size=8,
27     per_device_eval_batch_size=8,
28     num_train_epochs=3,
29     weight_decay=0.01,
30     save_total_limit=2,
31     logging_dir='./logs',
32     logging_steps=10,
33 )
34
35 # Define Trainer
36 trainer = Trainer(
37     model=model,
38     args=training_args,
39     train_dataset=train_dataset,
40     eval_dataset=val_dataset,
41 )
42 import os
43 os.environ["WANDB_API_KEY"] = "your_40_character_api_key"
44

```

```
46 # Train the model
47 trainer.train()
48
```

Some weights of BertForSequenceClassification were not initialized from the model chech You should probably TRAIN this model on a down-stream task to be able to use it for pr

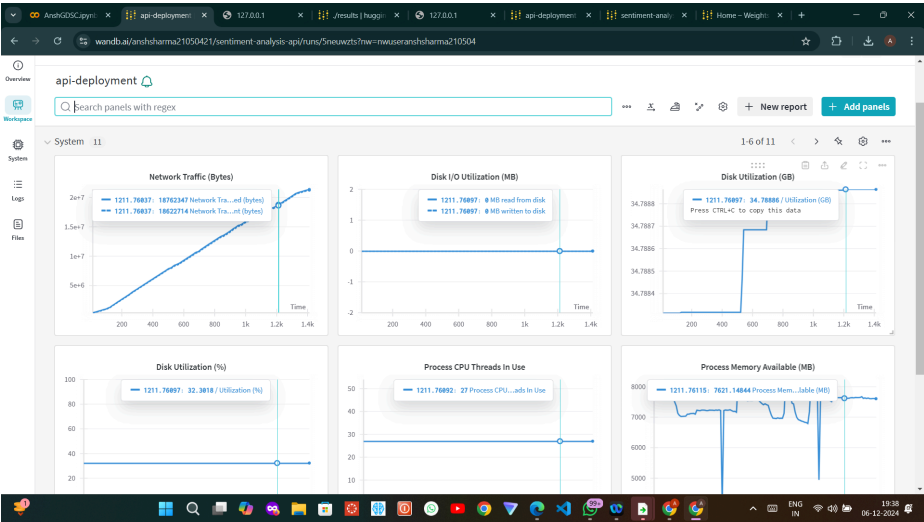
Map: 100% 3/3 [00:00<00:00, 27.68 examples/s]

Map: 100% 1/1 [00:00<00:00, 13.65 examples/s]

Map: 100% 1/1 [00:00<00:00, 15.50 examples/s]

/usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1568: FutureWarr
warnings.warn(
Changes to your `wandb` environment variables will be ignored because your `wandb` session has
already started. For more information on how to modify your settings with `wandb.init()` arguments, please
refer to [the W&B docs](#).
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb>
wandb: You can find your API key in your browser here: <https://wandb.ai/authorize>
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
Tracking run with wandb version 0.18.7
Run data is saved locally in /content/wandb/run-20241206_133753-iyw8jpp
Syncing run [./results](#) to [Weights & Biases \(docs\)](#)
View project at <https://wandb.ai/anshsharma21050421/huggingface>
View run at <https://wandb.ai/anshsharma21050421/huggingface/runs/iyw8jpp>
[3/3 01:06, Epoch 3/3]

Epoch	Training Loss	Validation Loss
1	No log	0.892518
2	No log	0.919221
3	No log	0.926235



```

1 # Evaluate the model on the test dataset
2 results = trainer.evaluate(test_dataset)
3
4 # Print evaluation results
5 print(f"Test Results: {results}")
6
7

```



[1/1 : <:]

✓ Evaluation of Model

```

1 from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
2
3 # Get predictions and true labels
4 predictions = trainer.predict(test_dataset)
5 preds = torch.argmax(torch.tensor(predictions.predictions), axis=1).numpy()
6
7 # Calculate accuracy and F1 score
8 accuracy = accuracy_score(test_dataset['label'], preds)
9 f1 = f1_score(test_dataset['label'], preds, average="weighted")
10
11 # Print the metrics
12 print(f"Accuracy: {accuracy}")
13 print(f"F1 Score: {f1}")
14
15 # Confusion Matrix
16 cm = confusion_matrix(test_dataset['label'], preds)
17 print("Confusion Matrix:")
18 print(cm)
19

```



Accuracy: 1.0
F1 Score: 1.0
Confusion Matrix:
[[1]]
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:409: UserWarning:



```

1 model.save_pretrained("./fine_tuned_bert")
2 tokenizer.save_pretrained("./fine_tuned_bert")
3

```



```

('./fine_tuned_bert/tokenizer_config.json',
 './fine_tuned_bert/special_tokens_map.json',
 './fine_tuned_bert/vocab.txt',
 './fine_tuned_bert/added_tokens.json')

```

```

1 from transformers import BertForSequenceClassification, BertTokenizer
2
3 # Load the fine-tuned model and tokenizer
4 model = BertForSequenceClassification.from_pretrained("./fine_tuned_bert")

```

```
5 tokenizer = BertTokenizer.from_pretrained("./fine_tuned_bert")
6
```

✓ Evaluation of Model by using Negative sentences and checking whether mode is predicting correct **sentiment**

```
1 # Example text
2 text = "Let's bomb NIT Calicut"
3
4 # Tokenize and prepare the input
5 inputs = tokenizer(text, return_tensors="pt", truncation=True, padding=True, max_length=6)
6
7 # Make prediction
8 with torch.no_grad():
9     outputs = model(**inputs)
10
11 # Get predicted sentiment
12 predicted_class = torch.argmax(outputs.logits, dim=1).item()
13
14 # Map the class index to sentiment
15 sentiment = {0: "Negative", 1: "Neutral", 2: "Positive"}
16 print(f"Sentiment: {sentiment[predicted_class]}")
17
```

➡ Sentiment: Negative

✓ FAST API

```
1 pip install fastapi uvicorn transformers torch wandb
```

➡ Collecting fastapi
 Downloading fastapi-0.115.6-py3-none-any.whl.metadata (27 kB)
 Collecting uvicorn
 Downloading uvicorn-0.32.1-py3-none-any.whl.metadata (6.6 kB)
 Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.8.0)
 Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.0.1)
 Requirement already satisfied: wandb in /usr/local/lib/python3.10/dist-packages (0.15.1)
 Collecting starlette<0.42.0,>=0.40.0 (from fastapi)
 Downloading starlette-0.41.3-py3-none-any.whl.metadata (6.0 kB)
 Requirement already satisfied: pydantic!=1.8,!=1.8.1,!=2.0.0,!=2.0.1,!=2.1.0,<3.0.0 in /usr/local/lib/python3.10/dist-packages (from starlette) (2.10.0)
 Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from pydantic) (4.12.2)
 Requirement already satisfied: click>=7.0 in /usr/local/lib/python3.10/dist-packages (from uvicorn) (8.1.8)
 Requirement already satisfied: h11>=0.8 in /usr/local/lib/python3.10/dist-packages (from uvicorn) (0.14.0)
 Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.16.1)
 Requirement already satisfied: huggingface-hub<1.0,>=0.23.2 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.23.2)
 Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.26.4)
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (24.1)
 Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.2)
 Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2024.11.6)
 Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from wandb) (2.32.0)
 Requirement already satisfied: tokenizers<0.21,>=0.20 in /usr/local/lib/python3.10/dist-packages (from wandb) (0.20.3)

```

Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-package
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist
Requirement already satisfied: docker-pycreds>=0.4.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: gitpython!=3.1.29,>=1.0.0 in /usr/local/lib/python3.
Requirement already satisfied: platformdirs in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: protobuf!=4.21.0,!5.28.0,<6,>=3.19.0 in /usr/local/
Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: sentry-sdk>=2.0.0 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: setproctitle in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-package
Requirement already satisfied: six>=1.4.0 in /usr/local/lib/python3.10/dist-package
Requirement already satisfied: gitdb<5,>=4.0.1 in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.10/
Requirement already satisfied: pydantic-core==2.27.1 in /usr/local/lib/python3.10/d
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.1
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist
Requirement already satisfied: anyio<5,>=3.4.0 in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-pa
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: exceptiongroup in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: smmap<6,>=3.0.1 in /usr/local/lib/python3.10/dist-pa
Downloading fastapi-0.115.6-py3-none-any.whl (94 kB)
    94.8/94.8 KB 3.4 MB/s eta 0:00:00
Downloading uvicorn-0.32.1-py3-none-any.whl (63 kB)
    63.8/63.8 KB 3.7 MB/s eta 0:00:00
Downloading starlette-0.41.3-py3-none-any.whl (73 kB)
    73.2/73.2 KB 4.4 MB/s eta 0:00:00
Installing collected packages: uvicorn, starlette, fastapi
Successfully installed fastapi-0.115.6 starlette-0.41.3 uvicorn-0.32.1

```

```

1 from fastapi import FastAPI
2 from pydantic import BaseModel
3 from transformers import BertForSequenceClassification, BertTokenizer
4 import torch
5 import wandb
6
7 # Initialize the FastAPI app
8 app = FastAPI()
9
10 # Load the fine-tuned model and tokenizer (replace with your saved model path)
11 model = BertForSequenceClassification.from_pretrained("./fine_tuned_bert")
12 tokenizer = BertTokenizer.from_pretrained("./fine_tuned_bert")
13
14 # Set up wandb (optional for experiment tracking)
15 wandb.init(project="sentiment-analysis-api", name="api-deployment")
16
17 # Define the input format using Pydantic
18 class Review(BaseModel):
19     text: str
20
21 # Define the sentiment prediction endpoint

```

```
22 @app.post("/predict/")
23 def predict_sentiment(review: Review):
24     # Tokenize input text
25     inputs = tokenizer(review.text, return_tensors="pt", truncation=True, padding=True,
26
27     # Make prediction
28     with torch.no_grad():
29         outputs = model(**inputs)
30
31     # Get the predicted sentiment (choose the max logit)
32     predicted_class = torch.argmax(outputs.logits, dim=1).item()
33
34     # Map class index to sentiment label
35     sentiment = {0: "Negative", 1: "Neutral", 2: "Positive"}
36     predicted_sentiment = sentiment[predicted_class]
37
38     # Log the prediction to wandb (optional)
39     wandb.log({"review": review.text, "predicted_sentiment": predicted_sentiment})
40
41     return {"review": review.text, "sentiment": predicted_sentiment}
42
```

↔ Changes to your `wandb` environment variables will be ignored because your `wandb` session has already started. For more information on how to modify your settings with `wandb.init()` arguments, please refer to [the W&B docs](#).
Finishing last run (ID:iynw8jpp) before initializing another...

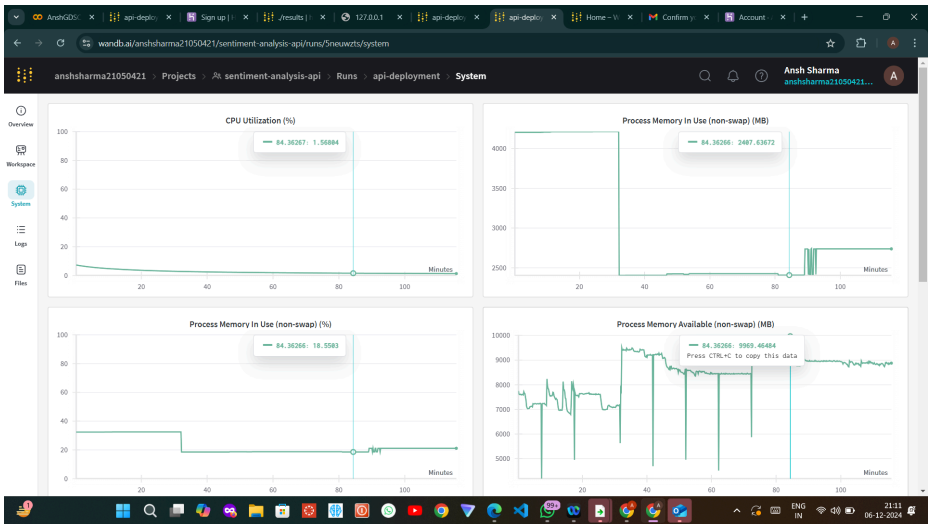
Run history:

eval/loss	
eval/runtime	
eval/samples_per_second	
eval/steps_per_second	
test/loss	
test/runtime	
test/samples_per_second	
test/steps_per_second	
train/epoch	
train/global_step	

Run summary:

eval/loss	1.02003
eval/runtime	3.4203
eval/samples_per_second	0.292
eval/steps_per_second	0.292
test/loss	1.02003
test/runtime	2.8443
test/samples_per_second	0.352
test/steps_per_second	0.352
total_flos	2368020759552.0
train/epoch	3
train/global_step	3
train_loss	1.08066
train_runtime	212.1883
train_samples_per_second	0.042
train_steps_per_second	0.014

View run **.results** at: <https://wandb.ai/anshsharma21050421/huggingface/runs/iynw8jpp>
View project at: <https://wandb.ai/anshsharma21050421/huggingface>
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)
Find logs at: `./wandb/run-20241206_133753-iynw8jpp/logs`
Successfully finished last run (ID:iynw8jpp). Initializing new run:
Tracking run with wandb version 0.18.7
Run data is saved locally in `/content/wandb/run-20241206_134511-5neuwztz`
Syncing run [api-deployment](#) to [Weights & Biases \(docs\)](#)



```
1 !curl https://cli-assets.heroku.com/install.sh | sh
2
```

```

% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload    Total     Spent    Left   Speed
100 1980    100 1980    0     0  6701      0 --:--:-- --:--:-- --:--:-- 6711
Installing CLI from https://cli-assets.heroku.com/channels/stable/heroku-linux-x64.tar
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload    Total     Spent    Left   Speed
100 51.4M    100 51.4M    0     0 3716k      0 0:00:14 0:00:14 --:--:-- 3279k
v16.20.2
heroku installed to /usr/local/bin/heroku
> Warning: Our terms of service have changed:
> https://dashboard.heroku.com/terms-of-service
heroku/9.5.1 linux-x64 node-v16.20.2

```

```
1 !heroku login -i
```

```

heroku: Enter your login credentials
Email: anshsharma210504@gmail.com
Password: *****

```

```
1 !pip freeze > requirements.txt
2
```

```

1 from fastapi import FastAPI
2 from pydantic import BaseModel
3 from transformers import BertForSequenceClassification, BertTokenizer
4 import torch
5 import wandb
6

```




```

7 app = FastAPI()
8
9 model = BertForSequenceClassification.from_pretrained("./fine_tuned_bert")
10 tokenizer = BertTokenizer.from_pretrained("./fine_tuned_bert")
11
12 class Review(BaseModel):
13     text: str
14
15 @app.post("/predict/")
16 def predict_sentiment(review: Review):
17     inputs = tokenizer(review.text, return_tensors="pt", truncation=True, padding=True,
18     with torch.no_grad():
19         outputs = model(**inputs)
20     predicted_class = torch.argmax(outputs.logits, dim=1).item()
21     sentiment = {0: "Negative", 1: "Neutral", 2: "Positive"}
22     predicted_sentiment = sentiment[predicted_class]
23     wandb.log({"review": review.text, "predicted_sentiment": predicted_sentiment})
24     return {"review": review.text, "sentiment": predicted_sentiment}
25

```

✓ Pushing GIT Repository and *Heroku* Deployment

```
1 !git init
```

 hint: Using 'master' as the name for the initial branch. This default branch name
 hint: is subject to change. To configure the initial branch name to use in all
 hint: of your new repositories, which will suppress this warning, call:
 hint:
 hint: git config --global init.defaultBranch <name>
 hint:
 hint: Names commonly chosen instead of 'master' are 'main', 'trunk' and
 hint: 'development'. The just-created branch can be renamed via this command:
 hint:
 hint: git branch -m <name>
 Initialized empty Git repository in /content/.git/

```


1 !git config --global user.email "ansh_b230825mt@nitc.ac.in"
2 !git config --global user.name "Anshsharmacse"

```

```

1 !git add .config/ fine_tuned_bert/ logs/ requirements.txt results/ sample_data/ wandb/
2

```

 .config/
 fine_tuned_bert
 logs
 requirements.txt
 results
 sample_data

```

1 !git commit -m "Initial commit"
2

```

```

[master (root-commit) 13426ee] Initial commit
56 files changed, 84591 insertions(+)
create mode 100644 .config/.last_opt_in_prompt.yaml
create mode 100644 .config/.last_survey_prompt.yaml
create mode 100644 .config/.last_update_check.json
create mode 100644 .config/active_config
create mode 100644 .config/config_sentinel
create mode 100644 .config/configurations/config_default
create mode 100644 .config/default_configs.db
create mode 100644 .config/gce
create mode 100644 .config/hidden_gcloud_config_universe_descriptor_data_cache_con
create mode 100644 .config/logs/2024.12.04/14.22.43.568001.log
create mode 100644 .config/logs/2024.12.04/14.23.06.536794.log
create mode 100644 .config/logs/2024.12.04/14.23.17.977201.log
create mode 100644 .config/logs/2024.12.04/14.23.19.373946.log
create mode 100644 .config/logs/2024.12.04/14.23.31.005371.log
create mode 100644 .config/logs/2024.12.04/14.23.31.672215.log
create mode 100644 fine_tuned_bert/config.json
create mode 100644 fine_tuned_bert/model.safetensors
create mode 100644 fine_tuned_bert/special_tokens_map.json
create mode 100644 fine_tuned_bert/tokenizer_config.json
create mode 100644 fine_tuned_bert/vocab.txt
create mode 100644 logs/events.out.tfevents.1733491855.223463f26f82.1619.0
create mode 100644 logs/events.out.tfevents.1733492156.223463f26f82.1619.1
create mode 100644 logs/events.out.tfevents.1733492454.223463f26f82.1619.2
create mode 100644 requirements.txt
create mode 100644 results/checkpoint-3/config.json
create mode 100644 results/checkpoint-3/model.safetensors
create mode 100644 results/checkpoint-3/optimizer.pt
create mode 100644 results/checkpoint-3/rng_state.pth
create mode 100644 results/checkpoint-3/scheduler.pt
create mode 100644 results/checkpoint-3/trainer_state.json
create mode 100644 results/checkpoint-3/training_args.bin
create mode 100755 sample_data/README.md
create mode 100755 sample_data/anscombe.json
create mode 100644 sample_data/california_housing_test.csv
create mode 100644 sample_data/california_housing_train.csv
create mode 100644 sample_data/mnist_test.csv
create mode 100644 sample_data/mnist_train_small.csv
create mode 120000 wandb/debug-internal.log
create mode 120000 wandb/debug.log
create mode 120000 wandb/latest-run
create mode 100644 wandb/run-20241206_133753-iyw8jpp/files/config.yaml
create mode 100644 wandb/run-20241206_133753-iyw8jpp/files/output.log
create mode 100644 wandb/run-20241206_133753-iyw8jpp/files/requirements.txt
create mode 100644 wandb/run-20241206_133753-iyw8jpp/files/wandb-metadata.json
create mode 100644 wandb/run-20241206_133753-iyw8jpp/files/wandb-summary.json
create mode 120000 wandb/run-20241206_133753-iyw8jpp/logs/debug-core.log
create mode 100644 wandb/run-20241206_133753-iyw8jpp/logs/debug-internal.log
create mode 100644 wandb/run-20241206_133753-iyw8jpp/logs/debug.log
create mode 100644 wandb/run-20241206_133753-iyw8jpp/run-iyw8jpp.wandb
create mode 100644 wandb/run-20241206_134511-5neuwzts/files/output.log
create mode 100644 wandb/run-20241206_134511-5neuwzts/files/requirements.txt
create mode 100644 wandb/run-20241206_134511-5neuwzts/files/wandb-metadata.json
create mode 120000 wandb/run-20241206_134511-5neuwzts/logs/debug-core.log
create mode 100644 wandb/run-20241206_134511-5neuwzts/logs/debug-internal.log
create mode 100644 wandb/run-20241206_134511-5neuwzts/logs/debug.log

```

1 !heroku create <Ansh Sentimental Analysis>

```

↵ /bin/bash: -c: line 1: syntax error near unexpected token `newline'
/bin/bash: -c: line 1: `heroku create <Ansh Sentimental Analysis>'

```

```
1 !heroku apps
```

```
↵ heroku: Press any key to open up the browser to login or q to exit:
```

```
1 !heroku create Anshsharmacse
```

```

↵ Creating ● Anshsharmacse...
Creating ● Anshsharmacse... ❖
Creating ● Anshsharmacse... ?
heroku: Press any key to open up the browser to login or q to exit: Creating ● Anshsh

```

```
1 !git remote add heroku https://git.heroku.com/Anshsharmacse.git
```

```
↵ error: remote heroku already exists.
```

```
1 !git remote -v
2
```

```

↵ heroku https://git.heroku.com/your-app-name.git (fetch)
heroku https://git.heroku.com/your-app-name.git (push)

```

```
1 !heroku login -i
```

```

↵ heroku: Enter your login credentials
Email: anshsharma210504@gmail.com
Password: *****

```

✓ Final Evaluation

Steps

1. Write the Review Text in the review box
2. Click on Predict Sentiment

Note→Refresh the cell on every new sentiment , to obtain the accurate result

```

1 # Import libraries
2 from transformers import BertForSequenceClassification, BertTokenizer
3 import torch
4 import ipywidgets as widgets

```

```
5 from IPython.display import display
6
7 # Load your fine-tuned model and tokenizer (replace with your model path if necessary)
8 model = BertForSequenceClassification.from_pretrained("./fine_tuned_bert")
9 tokenizer = BertTokenizer.from_pretrained("./fine_tuned_bert")
10
11 # Create a text box for user input
12 input_text = widgets.Textarea(
13     value='',
14     placeholder='Type your review here...',
15     description='Review Text:',
16     disabled=False,
17     layout=widgets.Layout(width='50%', height='100px')
18 )
19
20 # Define a function to predict sentiment
21 def predict_sentiment(text):
22     # Tokenize input text
23     inputs = tokenizer(text, return_tensors="pt", truncation=True, padding=True, max_le
24
25     # Make prediction
26     with torch.no_grad():
27         outputs = model(**inputs)
28
29     # Get predicted sentiment (choose the max logit)
30     predicted_class = torch.argmax(outputs.logits, dim=1).item()
31
32     # Map class index to sentiment label
33     sentiment = {0: "Negative", 1: "Neutral", 2: "Positive"}
34     predicted_sentiment = sentiment[predicted_class]
35
36     # Display the result
37     print(f"Predicted Sentiment: {predicted_sentiment}")
38
39 # Create a button to trigger the prediction
40 predict_button = widgets.Button(description="Predict Sentiment")
41
42 # Link the button to the prediction function
43 predict_button.on_click(lambda x: predict_sentiment(input_text.value))
44
45 # Display the input field and the button
46 display(input_text, predict_button)
47
```



Review Text: GDSC is the best club , Ansh
Sharma will certainly be
selected in GDSC Club.

Predict Sentiment

Predicted Sentiment: Positive

